

Martin Bergstø

Ontology-Driven Decision Support for Machine Learning Method Selection in Product Development and Production

Combining Semantic Literature Retrieval with
Ontology-Based Filtering and Implementation Support

Master's thesis in Engineering and ICT

Supervisor: Andrei Lobov

June 2026

Norwegian University of Science and Technology

Faculty of Engineering

Department of Mechanical and Industrial Engineering



ABSTRACT

Write an abstract/summary of your thesis, and state your main findings here.

A summary should be included in both English and any second language, if this is applicable, regardless if the thesis is written in English or in your preferred language. These should be on separate pages, the English version first.

PREFACE

Write the preface of your thesis here.

You may include acknowledgements and thanks as part of your preface on this page, or you may add it as a new chapter after the preface.

CONTENTS

Abstract	i
Preface	ii
Contents	iv
List of Figures	iv
List of Tables	v
Abbreviations	vii
1 Introduction	1
1.1 Motivation	1
1.2 Problem Description	2
1.3 Research Questions	3
1.4 Thesis Structure	3
2 Background and State of the Art	5
2.1 Machine Learning	5
2.1.1 Machine Learning Paradigms and Tasks	5
2.1.2 Approaches to Method Selection	6
2.2 Machine Learning in Product Development and Production	7
2.2.1 Applications of Machine Learning	8
2.2.2 Learning Paradigms in Industrial Applications	8
2.2.3 Pre-Master Findings	8
2.3 Ontologies and Knowledge Representation	9
2.4 Semantic Retrieval and Text Similarity	9
2.5 Related Work and Existing Tools	9
3 Methodology	11
4 Implementation	13
5 Results	15
6 Discussion	17

7	Conclusions	19
	References	21
	Appendices:	25

LIST OF FIGURES

2.1.1 Overview of machine learning paradigms with common tasks.	6
---	---

LIST OF TABLES

ABBREVIATIONS

List of all abbreviations in alphabetic order:

- **NTNU** Norwegian University of Science and Technology

INTRODUCTION

This chapter introduces the purpose and overall direction of the thesis. It presents the motivation for the study and describes the main problem that is addressed. The research questions are then defined to guide the work. Finally, the chapter provides an overview of the structure of the thesis.

1.1 Motivation

Machine learning has seen strong growth in recent years and is now widely applied in product development and production. It is used for tasks such as optimization, control, prediction, and decision support [1].

Advances in data availability and computing power have further increased the use of machine learning in industrial settings [2]. As a result, machine learning plays an important role in improving efficiency, supporting decision-making, and enabling new capabilities in engineering and manufacturing.

At the same time, the field of machine learning is large and diverse. Many different algorithms, methods, and approaches are available, each with its own strengths and weaknesses [1]. There is no single method that performs best for all problems, and performance depends on the specific application and data [3, 4].

This creates a challenge for engineers and practitioners, who must select suitable methods for their problems. The large number of available options can act as a barrier and limit the effective use of machine learning in practice [3]. Selecting an appropriate method often requires balancing multiple factors such as accuracy, interpretability, and data requirements [5, 6].

If the selection is not done carefully, the resulting model may perform poorly or fail to provide useful insights [6]. In practice, simpler models are often preferred due to their interpretability, even if more complex models may offer higher accuracy [7]. This highlights the importance of considering trade-offs in real-world applications.

Another challenge is the gap between research and practical implementation. Research results are distributed across many articles and domains, which makes it difficult to gain a complete overview [3]. In addition, most studies focus on developing models, while less attention is given to how they should be selected and applied in practice.

Companies also face challenges related to lack of knowledge and experience when adopting machine learning [8]. In many cases, the selection of methods is based on trial and error, which is time-consuming and inefficient [8]. There is also often pressure to adopt machine learning without a clear plan, which increases the risk of unsuccessful implementations [5].

These challenges highlight the need for more structured approaches that can support method selection and make machine learning more accessible.

1.2 Problem Description

Despite the growing adoption of machine learning in manufacturing and engineering, practitioners often lack adequate support when selecting appropriate methods for their specific problems. Although a large number of use cases exist, there is currently no guidance or standard for developing machine learning solutions from ideation to deployment, and it remains difficult for non-experts to begin developing solutions for their specific problems [9]. In practice, method selection is often based on trial and error, which is both time-consuming and dependent on prior experience that many practitioners do not have [8].

A central limitation is that no structured approach currently exists for making an informed selection of machine learning algorithms based on a well-defined problem description [8]. Existing frameworks tend to use accuracy or prediction speed as the primary selection criteria, which is often too general and can lead to unsuitable methods being chosen [10]. Manual evaluation of candidate algorithms is highly time-consuming, and if the selection is not done carefully, the resulting model may fail to produce useful results [6].

Existing research further focuses primarily on the development and evaluation of machine learning models, while less attention is given to what conditions lead to selecting one approach over another in specific industrial contexts [11]. The lack of interpretability in many approaches is also a key barrier to adoption, as practitioners need to understand and trust the recommendations they receive [12].

At the same time, the relevant knowledge is distributed across a large and fragmented body of literature. The extensive number of works on machine learning algorithms and applications creates broad knowledge on the topic, but may also generate disorientation when selecting the right approach [10]. This knowledge is difficult to access and apply systematically during method selection, which limits the ability to make informed, evidence-based decisions [3].

Addressing these challenges requires a structured approach that can translate problem descriptions into method recommendations and connect those recommendations to relevant research. This thesis investigates whether combining ontology-

based knowledge representation with semantic retrieval from scientific literature can provide such support, offering interpretable recommendations grounded in prior research alongside guidance for early-stage implementation.

Burde MLguide forklares i så stor detalj? Og bør det forklares her eller under research questions?

1.3 Research Questions

The challenges outlined above point to a need for structured decision support that can connect problem descriptions with suitable machine learning methods, use the research literature to justify those recommendations, and lower the barrier to early-stage implementation. This thesis investigates whether a system combining these capabilities can provide useful and justified recommendations in practice. The following hypothesis is proposed:

A decision-support system that combines ontology-based knowledge representation, semantic similarity between problem descriptions and scientific literature, and structured implementation support can provide relevant and justified machine learning method recommendations for users with limited machine learning expertise.

To investigate this hypothesis, the following research questions are defined:

- How can ontology-based knowledge improve machine learning method selection?
- How can semantic similarity between problem descriptions and literature support recommendations?
- How can literature-based evidence support and justify machine learning method recommendations?
- To what extent can such a system support early-stage machine learning implementation?

To investigate the hypothesis and research questions, a web-based decision-support system, called MLguide, is developed that combines structured knowledge with information from scientific literature.

1.4 Thesis Structure

The remainder of this thesis is structured as follows:

- Chapter 1 introduces the motivation, problem description, and research questions
- Chapter 2 presents background and related work
- Chapter 3 describes the methodology
- Chapter 4 presents the system implementation

- Chapter 5 presents the results
- Chapter 6 discusses the findings and limitations
- Chapter 7 concludes the thesis and suggests future work

BACKGROUND AND STATE OF THE ART

2.1 Machine Learning

- Overview of ML paradigms (supervised, unsupervised, etc.)
- Typical ML tasks (classification, regression, time series, etc.)
- Challenges in selecting appropriate methods
- Existing guidelines or frameworks for method selection

2.1.1 Machine Learning Paradigms and Tasks

Machine learning algorithms can be broadly categorized according to the type of supervision they receive during training. Mitchell [13] defines a learning algorithm as a computer program that improves its performance at some task through experience. Three main learning paradigms are commonly distinguished in the literature: supervised learning, unsupervised learning, and reinforcement learning [14] [15, Chapter 5] [16]. The following sections outline each paradigm and introduce some common tasks.

In supervised learning, the training data consists of input vectors paired with corresponding target values. The two most common supervised tasks are classification, where the goal is to assign an input to one of a finite number of discrete categories, and regression, where the goal is to predict a continuous numerical value [14] [15, Chapter 5].

In unsupervised learning, no target values are provided. The algorithm must instead identify structure in the data directly [14]. Common tasks include clustering, dimensionality reduction, and anomaly detection [16] [15, Chapter 5]. Clustering groups data points such that objects within the same group are more similar to each other than to those in other groups [16]. Dimensionality reduction simplifies data by reducing the number of features, which leads to better human

interpretations and lower computational costs [16]. Anomaly detection identifies observations that deviate significantly from the expected distribution, and is used in applications such as fraud detection and fault monitoring [15, Chapter 5]. Semi-supervised learning occupies a middle ground, operating on both labeled and unlabeled data, which is useful in contexts where labeled examples are scarce [16]. It is worth noting that the boundary between supervised and unsupervised learning is not always well-defined, and many methods can be applied across both types of problems [15, Chapter 5].

Reinforcement learning differs fundamentally from the other paradigms. Rather than learning from a fixed dataset, an agent interacts with an environment and learns through trial and error, receiving rewards or penalties based on its actions [14] [16]. Within reinforcement learning, methods can be broadly divided according to what they learn. In value-based methods, a value function is learned from which the policy is derived either explicitly or implicitly, whereas in policy-based methods the policy is learned directly without the need for a value function [17] [18, Chapter 11]. A third class, actor-critic methods, combines both approaches by maintaining a separate policy structure and value function simultaneously [18, Chapter 11].




Figure 2.1.1: Overview of machine learning paradigms with common tasks.

Sett inn et delkapittel om ML metoder? F.eks: Formidle at antallet konkrete metoder er stort. Vise at metodene varierer langs noen få meningsfulle dimensjoner som faktisk har praktisk betydning: tolkbarhet, datakrav, robusthet, og ytelse

2.1.2 Approaches to Method Selection

Selecting an appropriate machine learning method is not straightforward. As discussed in Chapter 1.1, the large number of available algorithms, each with distinct strengths and limitations, creates a real challenge for practitioners. This challenge is formally described by the Algorithm Selection Problem [19], which asks which algorithm is likely to perform best for a given problem. The No Free Lunch theorems provide a theoretical basis for this difficulty, showing that no single algorithm performs best across all problems and that good selection requires understanding the characteristics of the problem at hand [20].

Several practical tools have been developed to help with method selection. Scikit-learn provides a publicly available decision flowchart that guides users through questions about data size and task type to suggest candidate algorithms [21], and

Microsoft Azure offers a similar cheat sheet for its machine learning platform [22]. While these tools are accessible and widely used, they have clear limitations. The selection criteria used in both tools are mainly accuracy and prediction speed, something that can in practice lead to poor results. [10]. Factors such as interpretability, data characteristics, and domain context are not considered, and neither do they offer any guidance on how the problem itself should be formulated. [9].

More structured frameworks have been proposed in the academic literature. Sala et al. [10] present a two-layer selection framework that incorporates algorithm characteristics such as scalability and robustness to outliers, in addition to learning type and response type. Domain-specific frameworks have also been proposed, such as that of Arena et al. [11], who develop a structured set of decision variables for algorithm selection in predictive maintenance. However, as the authors note, most existing approaches focus narrowly on accuracy and computational requirements, and none provides a tool that connects algorithmic choices to the broader context of the application domain.

A related class of approaches is automated machine learning (AutoML), which seeks to automatically select, compose, and parametrize machine learning models to achieve optimal performance on a given task or dataset [23]. AutoML systems automate parts of the machine learning pipeline, including data preparation, feature engineering, hyperparameter optimization, and model generation, with the goal of reducing the demand for expert knowledge and making machine learning more accessible to domain experts [24]. Within AutoML, the algorithm selection problem is typically treated as an optimization task, where dataset characteristics are used to predict which algorithm is likely to perform best [25].

However, AutoML approaches have notable limitations in the context of decision support for non-experts. Most systems treat the selection and configuration process as a black box, providing little explanation of why a particular configuration was chosen [25]. This lack of transparency is a significant barrier to adoption, as users who do not understand or trust the output of a system are unlikely to apply it in practice [23]. Furthermore, AutoML primarily addresses the later stages of the machine learning pipeline, and the phases considered inherently human, such as problem formulation and interpretation of results, have received little attention [25]. This means that AutoML does not address the earlier challenge of helping practitioners understand which type of problem they have and which general approach is appropriate. This gap motivates the work in this thesis.

Nevne spesifikke AutoML verktøy som er relevante? Eller ta det i exisiting tools

2.2 Machine Learning in Product Development and Production

- Go through findings from the literature review in pre-master project
- ML is used for tasks such as Quality prediction and defect detection, Process

optimization, Predictive maintenance, Time series forecasting

- Product development often has limited and uncertain data
- Production systems may have structured and continuous data
- Selecting suitable ML methods is challenging in these contexts

2.2.1 Applications of Machine Learning

Machine learning is widely applied in industrial product development and production, where it supports tasks such as process optimization, quality prediction, fault diagnosis, predictive maintenance, and anomaly detection [1, 26]. The application of ML in these domains continues to grow, driven by the increasing availability of production data and the improved usability of ML tools [1]. Unlike traditional rule-based approaches, ML methods can learn directly from data and adapt to changing conditions, making them well suited to the complex and dynamic nature of manufacturing environments [1]. In practice, ML is used both to improve product quality and to optimize production processes, for example through early prediction of manufacturing outcomes, root cause analysis, and condition monitoring of equipment [27]. As well as individual processes, ML also supports broader operational tasks such as production planning, scheduling, and supply chain management [26, 12].

2.2.2 Learning Paradigms in Industrial Applications

Across reported applications in manufacturing and production, supervised learning is the dominant paradigm, accounting for the large majority of studies in most application areas [1, 28]. This reflects the nature of many industrial problems, where labeled data are available through quality inspections, sensor logs, and operator feedback, and where the goal is typically to predict a known outcome such as product quality or equipment failure [1, 12]. Unsupervised learning plays a smaller but important role, particularly in anomaly detection and condition monitoring, where labeled examples of faults are scarce or costly to obtain [12]. Reinforcement learning is less commonly reported in industrial settings, though it has shown strong results in process control and scheduling problems, where decisions must be made sequentially over time [29]. In practice, the choice of learning paradigm is often shaped by what data is available rather than by the structure of the problem itself [1].

2.2.3 Pre-Master Findings

In a preceding specialization project, a large-scale analysis of machine learning research in product development and production was conducted, classifying over 33,000 article abstracts across product lifecycle phase, application area, learning paradigm, and reported ML methods. The results show that the literature is strongly concentrated in the optimization phase of the product lifecycle, while early phases such as planning and concept development contain comparatively few studies. Supervised learning dominates across all phases and application areas, accounting for approximately 73–80% of the articles in each lifecycle phase,

with neural networks, random forests, and support vector machines as the most frequently reported methods. Based on these findings, a roadmap was proposed that structures the literature across lifecycle phase, application area, and learning paradigm to support early-stage method selection [30].

2.3 Ontologies and Knowledge Representation

- What an ontology is (concepts, relations, structure)
- Use of ontologies in engineering and ML domains
- Benefits: structure, explainability, reasoning
- Limitations: manual effort, rigidity

2.4 Semantic Retrieval and Text Similarity

- Text can be represented as embeddings (vectors)
- Similarity measures can compare meaning between texts
- Used to match problem descriptions with relevant content
- More flexible than keyword-based search

Literature-Based Discovery / Retrieval

- Scientific articles are a rich but unstructured knowledge source
- Challenges include large volume, noise, and unclear relevance
- Methods can be linked to articles through extraction and metadata
- Semantic retrieval can improve matching between problems and literature

2.5 Related Work and Existing Tools

- AutoML: Auto-sklearn, TPOT, H2O AutoML, Google AutoML
- Meta-learning: algorithm selection based on dataset characteristics and prior experience
- OMA-ML: ontology-based ML method selection
- Traingenerator: template-based notebook generation

CHAPTER
THREE

METHODOLOGY

CHAPTER

FOUR

IMPLEMENTATION

CHAPTER

FIVE

RESULTS

CHAPTER

SIX

DISCUSSION

CHAPTER
SEVEN

CONCLUSIONS

REFERENCES

- [1] Thorsten Wuest et al. “Machine learning in manufacturing: advantages, challenges, and applications”. In: *Production & Manufacturing Research* 4.1 (2016). doi: 10.1080/21693277.2016.1192517, pp. 23–45. ISSN: null. DOI: 10.1080/21693277.2016.1192517. URL: <https://doi.org/10.1080/21693277.2016.1192517>.
- [2] Rahul Rai et al. “Machine learning in manufacturing and industry 4.0 applications”. In: *International Journal of Production Research* 59.16 (2021). doi: 10.1080/00207543.2021.1956675, pp. 4773–4778. ISSN: 0020-7543. DOI: 10.1080/00207543.2021.1956675. URL: <https://doi.org/10.1080/00207543.2021.1956675>.
- [3] Isaac Kofi Nti et al. “Applications of artificial intelligence in engineering and manufacturing: a systematic review”. In: *Journal of Intelligent Manufacturing* 33.6 (2022), pp. 1581–1601. ISSN: 1572-8145. DOI: 10.1007/s10845-021-01771-6. URL: <https://doi.org/10.1007/s10845-021-01771-6>.
- [4] In Lee and Yong Jae Shin. “Machine learning for enterprises: Applications, algorithm selection, and challenges”. In: *Business Horizons* 63.2 (2020), pp. 157–170. ISSN: 0007-6813. DOI: <https://doi.org/10.1016/j.bushor.2019.10.005>. URL: <https://www.sciencedirect.com/science/article/pii/S0007681319301521>.
- [5] Siby Jose Plathottam et al. “A review of artificial intelligence applications in manufacturing operations”. In: *Journal of Advanced Manufacturing and Processing* 5.3 (2023), e10159. DOI: <https://doi.org/10.1002/amp2.10159>. URL: <https://doi.org/10.1002/amp2.10159>.
- [6] Rahman Ali, Sungyoung Lee, and Tae Choong Chung. “Accurate multi-criteria decision making methodology for recommending machine learning algorithm”. In: *Expert Systems with Applications* 71 (2017), pp. 257–278. ISSN: 0957-4174. DOI: <https://doi.org/10.1016/j.eswa.2016.11.034>. URL: <https://www.sciencedirect.com/science/article/pii/S0957417416306698>.
- [7] Andrei Paleyes, Raoul-Gabriel Urma, and Neil D. Lawrence. “Challenges in Deploying Machine Learning: A Survey of Case Studies”. In: *ACM Comput. Surv.* 55.6 (2022), Article 114. ISSN: 0360-0300. DOI: 10.1145/3533378. URL: <https://doi.org/10.1145/3533378>.

- [8] Sebastian Sonntag et al. “A PATTERN LANGUAGE APPROACH TO IDENTIFY APPROPRIATE MACHINE LEARNING ALGORITHMS IN THE CONTEXT OF PRODUCT DEVELOPMENT”. In: *Proceedings of the Design Society 3* (2023), pp. 365–374. DOI: 10.1017/pds.2023.37. URL: <https://www.cambridge.org/core/product/44BBE29AA01344BC5F50174953C5BA8B>.
- [9] Tingting Chen et al. *Machine Learning in Manufacturing towards Industry 4.0: From ‘For Now’ to ‘Four-Know’*. Electronic Article. 2023. DOI: 10.3390/app13031903.
- [10] Roberto Sala et al. *How to select the right machine learning algorithm: a feature-based, scope-oriented selection framework*. 2018.
- [11] Simone Arena et al. “A conceptual framework for machine learning algorithm selection for predictive maintenance”. In: *Engineering Applications of Artificial Intelligence* 133 (2024), p. 108340. ISSN: 0952-1976. DOI: <https://doi.org/10.1016/j.engappai.2024.108340>. URL: <https://www.sciencedirect.com/science/article/pii/S0952197624004986>.
- [12] Anna Presciuttini et al. “Machine learning applications on IoT data in manufacturing operations and their interpretability implications: A systematic literature review”. In: *Journal of Manufacturing Systems* 74 (2024), pp. 477–486. ISSN: 0278-6125. DOI: <https://doi.org/10.1016/j.jmsy.2024.04.012>. URL: <https://www.sciencedirect.com/science/article/pii/S0278612524000773>.
- [13] Tom M. Mitchell. *Machine Learning*. New York: McGraw-Hill, 1997.
- [14] Christopher M. Bishop. *Pattern Recognition and Machine Learning*. Information Science and Statistics. New York: Springer, 2006. ISBN: 978-0387310732.
- [15] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. <http://www.deeplearningbook.org>. MIT Press, 2016.
- [16] Iqbal H. Sarker. “Machine Learning: Algorithms, Real-World Applications and Research Directions”. In: *SN Computer Science* 2.3 (2021), p. 160. ISSN: 2661-8907. DOI: 10.1007/s42979-021-00592-x. URL: <https://doi.org/10.1007/s42979-021-00592-x>.
- [17] Mohit Sewak. “Policy-Based Reinforcement Learning Approaches”. In: *Deep Reinforcement Learning: Frontiers of Artificial Intelligence*. Ed. by Mohit Sewak. Singapore: Springer Singapore, 2019, pp. 127–140. ISBN: 978-981-13-8285-7. DOI: 10.1007/978-981-13-8285-7_10. URL: https://doi.org/10.1007/978-981-13-8285-7_10.
- [18] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. 2nd. Cambridge, MA: MIT Press, 2018.
- [19] Kate A. Smith-Miles. “Cross-Disciplinary Perspectives on Meta-Learning for Algorithm Selection”. In: *ACM Computing Surveys* 41.1 (2009). DOI: 10.1145/1456650.1456656.
- [20] David H. Wolpert and William G. Macready. “No Free Lunch Theorems for Optimization”. In: *IEEE Transactions on Evolutionary Computation* 1.1 (1997), pp. 67–82. DOI: 10.1109/4235.585893.

- [21] Scikit-learn developers. *Choosing the right estimator*. Accessed: May 2026. URL: https://scikit-learn.org/stable/machine_learning_map.html.
- [22] Microsoft. *Machine Learning Algorithm Cheat Sheet for Azure Machine Learning*. Accessed: May 2026. URL: <https://learn.microsoft.com/en-us/azure/machine-learning/algorithm-cheat-sheet>.
- [23] Jonathan Waring, Charlotta Lindvall, and Renato Umeton. “Automated machine learning: Review of the state-of-the-art and opportunities for health-care”. In: *Artificial Intelligence in Medicine* 104 (2020), p. 101822. ISSN: 0933-3657. DOI: <https://doi.org/10.1016/j.artmed.2020.101822>. URL: <https://www.sciencedirect.com/science/article/pii/S0933365719310437>.
- [24] Xin He, Kaiyong Zhao, and Xiaowen Chu. “AutoML: A survey of the state-of-the-art”. In: *Knowledge-Based Systems* 212 (2021), p. 106622. ISSN: 0950-7051. DOI: <https://doi.org/10.1016/j.knosys.2020.106622>. URL: <https://www.sciencedirect.com/science/article/pii/S0950705120307516>.
- [25] Rafael Barbudo, Sebastián Ventura, and José Raúl Romero. “Eight years of AutoML: categorisation, review and trends”. In: *Knowledge and Information Systems* 65.12 (2023), pp. 5097–5149. ISSN: 0219-3116. DOI: 10.1007/s10115-023-01935-1. URL: <https://doi.org/10.1007/s10115-023-01935-1>.
- [26] Siby Jose Plathottam et al. “A review of artificial intelligence applications in manufacturing operations”. In: *Journal of Advanced Manufacturing and Processing* 5.3 (2023), e10159. DOI: <https://doi.org/10.1002/amp2.10159>. URL: <https://doi.org/10.1002/amp2.10159>.
- [27] Dorina Weichert et al. “A review of machine learning for the optimization of production processes”. In: *The International Journal of Advanced Manufacturing Technology* 104.5 (2019), pp. 1889–1902. ISSN: 1433-3015. DOI: 10.1007/s00170-019-03988-5. URL: <https://doi.org/10.1007/s00170-019-03988-5>.
- [28] Isaac Kofi Nti et al. “Applications of artificial intelligence in engineering and manufacturing: a systematic review”. In: *Journal of Intelligent Manufacturing* 33.6 (2022), pp. 1581–1601. ISSN: 1572-8145. DOI: 10.1007/s10845-021-01771-6. URL: <https://doi.org/10.1007/s10845-021-01771-6>.
- [29] Marcel Panzer and Benedict Bender. “Deep reinforcement learning in production systems: a systematic literature review”. In: *International Journal of Production Research* 60.13 (2022). doi: 10.1080/00207543.2021.1973138, pp. 4316–4341. ISSN: 0020-7543. DOI: 10.1080/00207543.2021.1973138. URL: <https://doi.org/10.1080/00207543.2021.1973138>.
- [30] Martin Bergstø. “Supporting Method Selection for Machine Learning in Product Development and Production: A Multi-Dimensional Classification of Machine Learning Research”. In: (Jan. 2026). Specialization Project.

APPENDICES